

1

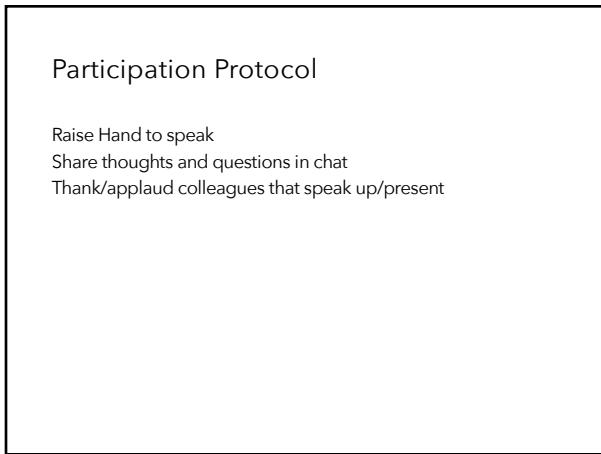
Using Zoom

1. Open Chat & Participants
2. Merge Window

Click the green check in the participants list when done

The screenshot shows a Zoom meeting window. It includes the "Participants" section and a "Chat" section. A red box highlights the "Merge Into... Window" button in the participant list. Red arrows point to the "Merge Into... Window" button and the "Participants" section. The bottom bar has controls like "Join Audio", "Share Screen", "Invite Others", "Invite", "Mute Me", and "Leave".

2



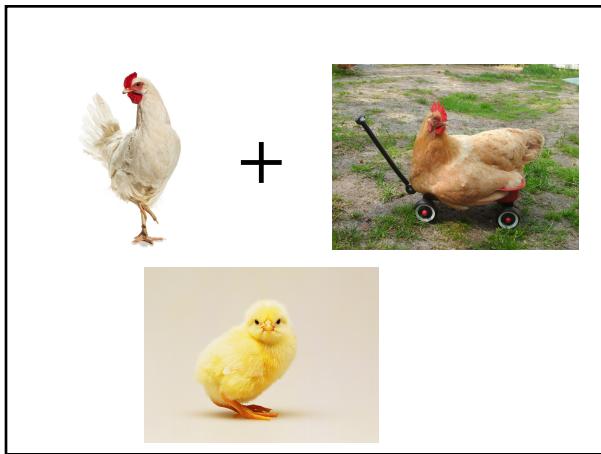
3

Is it possible for  $1 + 1 = 3$ ? If so, when?

Open chat and share what you think

The screenshot shows a Zoom meeting window with a "Chat" section at the bottom. A red box highlights the "Chat" section. The bottom bar has controls like "Join Audio", "Share Screen", "Invite Others", "Invite", "Mute Me", and "Leave".

4



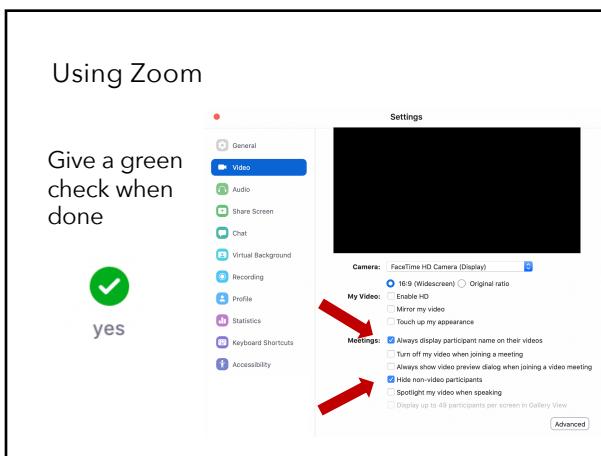
5

Participation Protocol

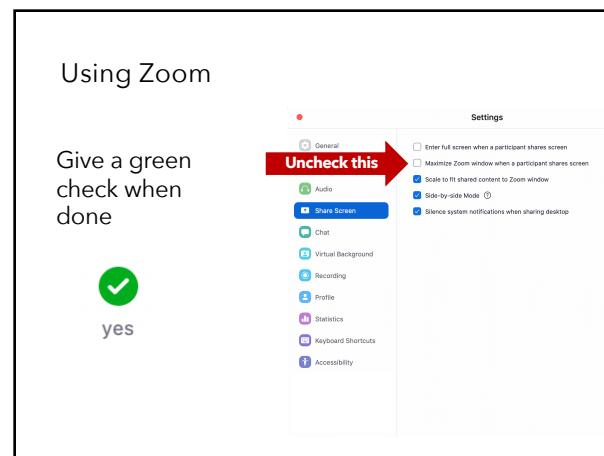
Raise Hand to speak  
Share thoughts and questions in chat  
Thank/appaud colleagues that speak up/present

I may ask individuals to share thoughts on the readings

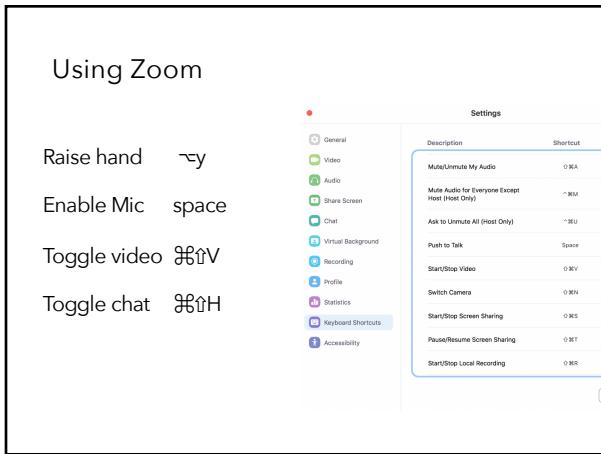
6



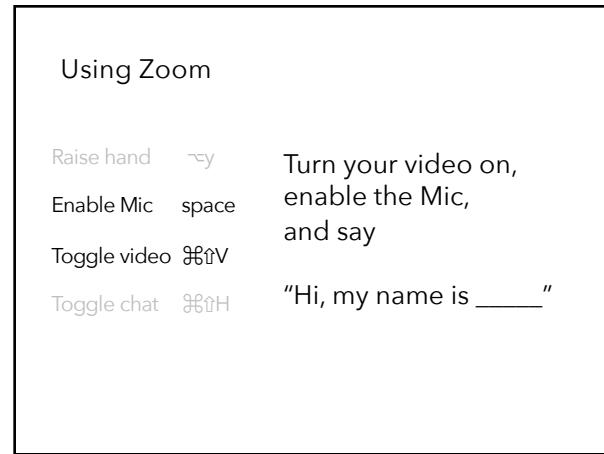
7



8



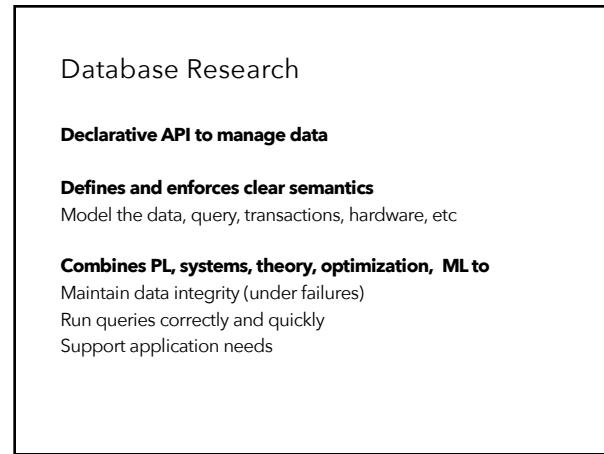
9



10



11



12

### Database Research

**Definition of data changing**

- Records
- Graphs
- Sensor events
- Images & Videos
- Documents & PDFs
- VR/AR...

**Application needs are growing**

- Hardware, scalability, ML, real-time
- Decision making, data exploration, video analytics, etc

### This Course

**Learn about Database Research**

- Read classic and modern papers
- Work on a research project

**Learn how to read papers**

- Problem selection
- Insights and technical contributions
- Presentation

13

14

### Course Topics

- Query Execution
- Query Optimization
- Columnar Stores
- Query Compilation
- Large-scale Dataflow
- Materialized Views
- Datalog and Lineage
- Physical database optimization
- <Your interest here>

### Course Topics

We will NOT cover material in 4111, namely

- Relational Algebra
- SQL
- Basics of query execution and optimization
- Access methods
- Basic design of disk-based database systems

15

16

### What you will do

**Project (65%)**  
Groups of 2-3

**Assignments (15%)**  
See how pieces of a DB engine fit together

**Paper Reviews (10%)**  
Answer questions for each paper & submit review on wiki  
Lots of reading!

**Class discussion (10%)**

### Project

**2 Broad Categories**

- Reproduce and extend
- Research project

17

18

## Reproduce (& Extend)

### **Reproduce**

Read and deeply understand a data management topic  
~5 to 10+ papers, summarize & compare them  
Implement best/combined technique in DataBass/another sys  
Benchmark it

### **& Extend (bonus)**

Extend the implementation to make it a bit better  
Benchmark to show it is better  
Explain when and why the extension is effective

## Research Project

### **Investigate new ideas and solutions**

Define hypothesis  
Conduct research  
Evaluate hypothesis  
Writeup and research

### **Emphasis on hypothesis and evaluation**

Translate contributions into testable hypotheses  
Evaluation should evaluate hypotheses

19

20

## Project Overview

Choose from list of projects, or come up with your own

### **Pick partners, 1 page proposal**

Problem you are solving  
Hypothesis  
Plan of attack, with milestones  
Preliminary related work

### **Demo/Presentation session**

### **Project report**

## Project Timeline

We are here to help you succeed

### Week

- |    |                                       |
|----|---------------------------------------|
| 02 | Release list of projects              |
| 04 | Submit proposal                       |
|    | Discuss with staff before this stage  |
| 07 | Submit paper draft                    |
|    | Position relative to state of the art |
|    | Play with related tools               |
| 11 | Check in                              |
| 14 | Showcase                              |
| 15 | Submit paper 8-10 pgs                 |

21

22

## Programming Assignments

**DataBass:** Python DB engine ~4K loc 

Hands-on experience with query from parsing to execution



- A0: add ORDERBY clause
- A1: hashjoin, pushdown optimization
- A2: selinger join optimization
- A3: hashjoin compilation
- A4: benchmarking
- A5: lineage

## Class Format

Submit paper review before class

Quiz for self assessment

Breakout to discuss quiz/reading

Go over paper ideas

Use chat to ask/answer questions

Optional: sign up to lead discussions

Optional: scribe the discussion, add to course wiki

23

24

## Reviews

Add to course wiki  
Can skip 4 **paper reviews** (note: 1+ papers per class)  
Late submissions not accepted

**Questions**

What is the problem?  
Why does prior work fail to solve it?  
Main insight and technical contributions?  
Do the evaluations hold up?  
What do you wish could be improved?

25

## Reviews

**Do not plagiarize**

Write original reviews  
Do **not** copy text from papers, online, or other reviews

See <http://www.cs.columbia.edu/education/honesty>

26

## Presentation

15-20 minutes + discussion

Cover **key elements** of paper  
Find & read related work to provide context  
Intuition >> formulae  
Create your own examples. Plagiarism rules apply

Lead discussion with questions

Go over presentation with Eugene ahead of time  
Send slides to TA 2 days before class, by midnight.

27

## Recitation

Tues 2-3PM EST, an open air space (113<sup>th</sup> & Morningside)  
Attendance is *not* taken



**Purpose**  
Go over examples  
Research advice  
Life advice  
*Stay safe from COVID*

28

## Who am I?

Eugene Wu  
421 Mudd (DSI space)  
ewu@cs.columbia.edu  
OH: Tues 3-4PM EST

PhD MIT, 2015  
Started @Columbia Fall 2015  
Systems for Human Data Interaction

29

## Deka Auliya Akbar

TA for the class  
OH: tba

30

## Logistics

### w6113.github.io

Class T/Th 11:30 – 1PM, Zoom  
 Recitation Tues 2-3PM, open air location TBD  
 Commun. Slack  
 Reviews Course wiki

### Academic Honesty

See <http://www.cs.columbia.edu/education/honesty>  
 Ask if unsure, falling behind, etc  
 Don't plagiarize. Worse than failing.

## Prereqs

### Required: 4111 Intro to DB

Declarative languages and data independence  
 Relational algebra, SQL, relational model  
 Query optimization and execution

### Useful: 4112

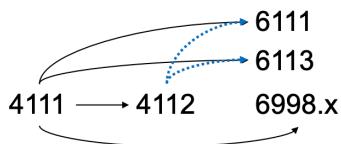
### Great to have people with different backgrounds

Talk to me if unsure  
 If graduate student from different area, talk to me

31

32

## Relationship with other DB classes



33

## Enrollment

~25 students  
 Doesn't matter if enrolled or waitlisted

Admission based on quality of  
 First reviews  
 Participation in initial classes

34

## Grading

Project	65%
Assignments	15%
Reviews	10%
Participation	10%

No curve, everyone can get A  
 If project is amazing, automatic A

## Breakout: Getting to know each other

Assign one note taker to summarize the breakout to the class  
 Introduce yourself  
 Your name, year, department  
 What is a database you recently used (maybe indirectly)?  
 Share where you want to travel when you are able to.

35

36

5 min Break

37

## Short DB History

### What Goes Around Comes Around

Michael Stonebraker  
Joseph M. Hellerstein

#### Abstract

This paper provides a summary of 14 years of database proposals presented in 9 different conferences. We discuss the overlaps of methods, and show that there are only a few basic data modeling ideas, and most have been around a long time. Late proposals inevitably bear a strong resemblance to certain earlier proposals. Hence, it is a good idea to review old proposals.

In addition, we present the lessons learned from the exploration of the proposals in each era. Most current researchers were not around for many of the previous eras, and have limited (if any) understanding of what was previously learned. There is an old adage that he who does not understand history is condemned to repeat it. By presenting "ancient"

38

## What is a Database?

### Data stored in a representation

Records and relationships

### Query the data

Write or generate algorithms to traverse storage representation

### Innovations found in

Data representation

Query interface

Translating query interface → algorithms over data

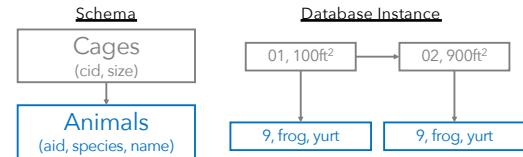
39

## 60s Hierarchical Model

### IMS: IBM Management System

For Apollo space program: Saturn V inventory  
Still used by banks today

Each level is a different entity type



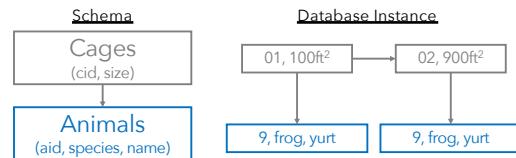
40

## 60s Hierarchical Model

```

# Animals living in cage 2
Find cage where cid = 2
Find 1st child of current record
Loop
    Find next sibling record of same type
    print record

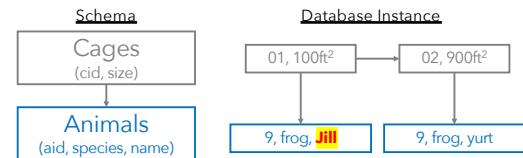
```



41

## 60s Hierarchical Model

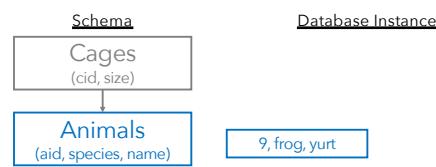
⚠ Data redundancy if many to many relationship



42

## 60s Hierarchical Model

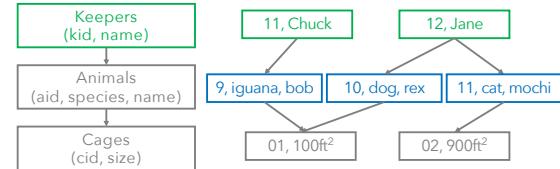
- Data redundancy if many to many relationship
- Animal can't exist without a cage
- Removing cage removes animals



43

## 60s Network Model (CODASYL)

- Charles Bachman Turing #8  
Record at a time navigation  
Impossible to program  
Changing data representation breaks programs

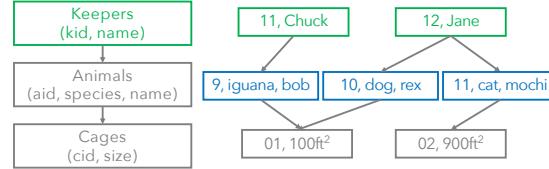


44

## 60s Network Model

```

Find keeper where name = Jane
Loop
    Find next animal in cares_for
    Find Cage in lives_in
    print record
  
```

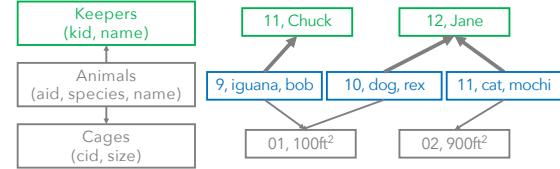


45

## 60s Network Model

```

Find keeper where name = Jane
Loop
    Find next animal in cares_for
    Find Cage in lives_in
    print record
  
```

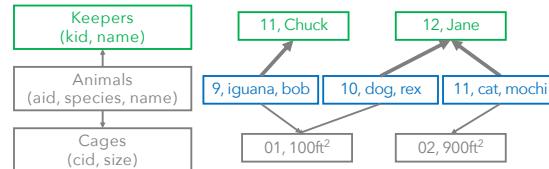


46

## 60s Network Model

```

Find keeper where name = Jane
Loop
    Find next animal in cares_for
    Find Cage in lives_in
    print record
  
```



47

## 70s Relational Model

Edgar Codd Turing #18  
Started the relational DB research field  
*Game changer*



Set at a time language  
Declarative language  
Data independence

**Information Retrieval**  
**A Relational Model of Data for Large Shared Data Banks**  
E.F. Codd  
IBM Research Laboratory, San Jose, California

The relational view of model of data described in Section I appears to be superior to several others in the literature in its ability to support the needs of large shared data banks. It provides a means of describing data which is well suited to the needs of large shared data banks and is well suited to the needs of users who will be interacting with such data banks through terminals or other devices for the purpose of information retrieval. It provides a basis for a high level programming language for manipulating data in large shared data banks. It provides a basis for writing programs on the one hand and reading results on the other hand. It provides a means of describing data which is well suited to the needs of large shared data banks and is well suited to the needs of users who will be interacting with such data banks through terminals or other devices for the purpose of information retrieval.

The relational view of model of data described in Section I appears to be superior to several others in the literature in its ability to support the needs of large shared data banks. It provides a means of describing data which is well suited to the needs of large shared data banks and is well suited to the needs of users who will be interacting with such data banks through terminals or other devices for the purpose of information retrieval. It provides a basis for a high level programming language for manipulating data in large shared data banks. It provides a basis for writing programs on the one hand and reading results on the other hand. It provides a means of describing data which is well suited to the needs of large shared data banks and is well suited to the needs of users who will be interacting with such data banks through terminals or other devices for the purpose of information retrieval.

The relational view of model of data described in Section I appears to be superior to several others in the literature in its ability to support the needs of large shared data banks. It provides a means of describing data which is well suited to the needs of large shared data banks and is well suited to the needs of users who will be interacting with such data banks through terminals or other devices for the purpose of information retrieval. It provides a basis for a high level programming language for manipulating data in large shared data banks. It provides a basis for writing programs on the one hand and reading results on the other hand. It provides a means of describing data which is well suited to the needs of large shared data banks and is well suited to the needs of users who will be interacting with such data banks through terminals or other devices for the purpose of information retrieval.

48

**70s Relational Model**

kid	name
11	Chuck
12	Jan

cid	size
1	100ft <sup>2</sup>
2	900ft <sup>2</sup>

kid	aid
11	9
12	10
12	11

cid	aid
1	9
1	10
2	11

aid	name	name
9	iguana	bob
10	dog	rex
11	cat	mochi

**70s Relational Model**

kid	name
11	Chuck
12	Jan

cid	size
1	100ft <sup>2</sup>
2	900ft <sup>2</sup>

kid	aid
11	9
12	10
12	11

cid	aid
1	9
1	10
2	11

aid	name	name
9	iguana	bob
10	dog	rex
11	cat	mochi

49

50

**70s Relational Model**

When do you want data independence?

δApp << δenvironment  
- Hellerstein

Application logic evolves slowly  
Technology evolves rapidly  
Change data layout  
Change schema  
Change/upgrade hardware  
Data system is remote  
Cloud  
...

**70s Will it run?**

**Ingres @Berkeley**  
Stonebraker Turing #32 and Wang  
Begat Ingres, Bitton-lee, Sybase, MS SQLServer, PACE, Tandem

**System R @IBM**  
Jim Gray Turing #22  
Big research team, 15 PhDs  
Begat DB2, Oracle, HP Allbase, Tandem

Proved that theory is practical  
Spawned RDBMS market and modern data-oriented software

51

52

**80s Will it Sell?**

**Heavyweights**  
IBM doesn't want to cannibalize IMS  
Oracle reads System R papers, and goes to market first  
IBM eventually releases DB2  
IBM chooses SQL, Oracle follows, becomes the standard

**Crowd of players**  
Jim Gray et al join Tandem  
Stonebraker makes Ingres → Informix  
Britton & Lee from Ingres create Britton-Lee  
Epstein leaves BL to create SyBase (bought by SAP for \$5B)  
Wang creates PACE  
Caltech alums start Teradata w/ networking tech → parallel DB

**80s Object Oriented DBMS**

Impedance mismatch between Objects and relational model

```
Struct person {
    name string;
    hobbies string[];
}
```

pId	name
1	Chuck
2	Jane

pId	hobby
1	cheese
1	surfing

Translation cumbersome  
Wanted to avoid the join  
Prefer DOT notation e.g., person.hobbies

53

54

## 80s Object Oriented DBMS

Impedance mismatch between Objects and relational model

```
Struct person {
    name string;
    hobbies string[];
}

person
pid | name | hobbies
---|---|---
1   | Chuck | [cheese,
           |         surfing]
2   | Jane  | [cheese,
           |         surfing]
```

```
{
    pid: 1
    name: "Chuck"
    hobbies: [
        "cheese",
        "surfing"
    ]
}
```

## Nested Relational Model

55

## 80s Object Oriented DBMS

Impedance mismatch between Objects and relational model

```
Struct person {
    name string;
    hobbies string[];
}

person
pid | name | hobbies
---|---|---
1   | Chuck | [cheese,
           |         surfing]
2   | Jane  | [cheese,
           |         surfing]
```

```
{
    pid: 1
    name: "Chuck"
    hobbies: [
        "cheese",
        "surfing"
    ]
}
```

Hard to query against  
 Vendor lock-in, language lock-in  
 NoSQL before it was cool  
 Mostly addressed by ORMs

56

## 90s More Money

*Sybase* licensed by Microsoft for x86 → *Microsoft SQLServer*

*MySQL* created

*Postgres* adds support for SQL → *PostgreSQL*

*SQLite* created.  
 Most popular DB in the world  
 Second most popular software library in the world (after zlib)

57

## 00s Data Warehouses aka OLAP

Companies accumulate transaction data, want to analyze  
 Want high throughput (OLAP)  
 Distributed shared nothing databases  
 Mostly closed source  
 Limited scalability (tens of nodes)

*PostgreSQL* begats *Netezza*, *Greenplum*  
*MonetDB* from CWI  
*Vertica* from MIT  
*Paraccel* → *RedShift*

58

## 00s The Internet

Internet companies want scalability over all else  
 Give up most of RDBMS features for scalability  
 Schema, relational model, ACID, SQL, strong consistency  
 Rise of eventually consistent NoSQL stores  
 Open Source  
*BigTable*, *Dynamo*, *Hbase*, *MongoDB*, *Couchbase*, *Cassandra*

59

## 10s NewSQL

Eventual consistency is difficult to program against  
 Scalable distributed DBs that support ACID, SQL  
 Shared nothing via partitioning (sharding)

*Industry*: *Google Spanner*, *MemSQL*, *SAP Hana*  
*Academic* → *Company*: *Hstore* → *VoltDB*, *Hyper*,  
*Open Source*: *Yugabyte*, *Cockroach*

60

## 10s Cloud Databases

Cloud infrastructure = Scalable resource provisioning

**Early on** containerize open source databases  
Provides devops, provisioning  
Doesn't best leverage cloud  
*AWS Postgresql, Amazon RDS*

**Cloud native databases**  
Optimized for distributed storage (shared disk)  
Decouples storage and execution, scales independently  
Lots of connectors to the "data lake"  
*Presto, Snowflake, Apache Drill, AWS Redshift, Azure Cloud*

## 10s Specialized DBs

**Graph systems**  
Graph API over relational (node, edge) data  
CIDR 2015 GRAL: fast DBMS > specialized graph engines  
Adage: graph traversals are joins. DBs are very good at joins  
*Neo4J, Dgraph, Graphbase, Giraph, GraphX*

**Time series**  
Timestamp + a few attributes  
Custom indexing, execution for high performance  
*Timescale, InfluxDB, clickhouse,*

61

62

## 10s Specialized DBs

**Streaming**  
Events arrive in a "stream"  
Querying over an infinitely growing table

**Evolution of streaming systems**  
Continuous queries: SASE, Telegraph, STREAM  
Sliding windows: Aurora, Oracle CQL  
Scalability, Best-effort semantics: Twitter Storm  
Out of order, state: Flink, Spark Streaming, Kafka, Materialized

## 10s DBs for X

Classic DB architecture are still good bones  
RDBMS is mostly commodity  
Market shift toward "massive scale" and "for X"

**Where X is**  
Ads  
Video  
Astronomy  
ML  
Vis  
...

**Rule of thumb**  
10x better not enough  
>50x better

63

64

At the end of the day

Main Market: DB **Apps & Services**

**Salesforce:** CRM  
**IBM, SAP, Oracle:** CRM, supply chain, inventory, HR  
**Adobe:** marketing, advertisement  
**Google, FB, Yandex, Twitter:** the web, search  
**Microsoft:** productivity, cloud  
**Amazon:** commerce, cloud

Each company created major DBMSes

At the end of the day

Main Market: DB **Apps & Services**

Database Management Technology

65

66

Your Tasks for next class

Submit reviews for Thursday

Submit Assignment 0 by 9/13 11:59PM EST

w6113.github.io